

INCIDENT RESPONSE PLAN

PayGuard Fintech (Illustrative MVP)

A Practical Framework for Detecting, Containing, and Recovering from Security Incidents

3MTT Capstone Project — Cybersecurity Incident Response

Document Version 1.0 | August 2026

Classification: Internal — Confidential

Table of Contents

1. Purpose and Scope	3
1.1 Objectives	3
1.2 Scope.....	3
1.3 Out of Scope	3
2. Roles and Responsibilities	3
2.1 Escalation Rule	4
2.2 Backup Contacts.....	4
3. Severity Classification	4
4. Incident Playbooks	5
4.1 Playbook A — Account Takeover / Credential Stuffing	5
4.2 Playbook B — Unauthorized Server / Infrastructure Access	6
4.3 Playbook C — Payment Fraud / Suspicious Transaction Spike	7
5. Communications Plan	8
5.1 Internal Communications	8
5.2 Customer Communications	8
5.3 Regulatory Communications.....	8
6. Reporting	9
6.1 Incident Timeline Template	9
6.2 Post-Incident Report Structure.....	10
7. Remediation & Hardening Recommendations.....	10
8. Appendix.....	10
8.1 Contact & Escalation Roster	10
8.2 Lab Environment Notes (for Assessment Evidence)	11
8.3 Tools Reference	11
8.4 Revision History	11
9. Glossary	11
10. References	13
10.1 Note on Citation Practice.....	14

1. Purpose and Scope

Nigerian fintech startups are high-value, high-frequency targets for account takeover, credential stuffing, payment fraud, and unauthorized system access — yet most early-stage teams have no documented response process. When an incident happens, panic replaces process: nobody knows who decides, who talks to customers, or what the regulator requires, and by the time someone thinks to write anything down, the timeline for CBN's 24-hour reporting window is already gone.

This plan gives PayGuard (used here as an illustrative early-stage Nigerian fintech) a lightweight, actionable incident response process that a small team can actually execute under pressure. It defines who does what, how to detect and contain the three incident types most relevant to a Nigerian payment's startup, how and when to communicate internally, to customers, and to regulators, and how to document everything so the resulting report is defensible.

1.1 Objectives

- Reduce time-to-containment by removing decision paralysis during an active incident.
- Meet Nigerian regulatory notification obligations (NDPC, CBN) without missing deadlines.
- Preserve evidence in a way that supports root-cause analysis and, if needed, law enforcement referral.
- Produce a consistent post-incident report that feeds back into hardening the platform.

1.2 Scope

This plan covers security incidents affecting PayGuard's production and staging environments, customer accounts, payment processing systems, and any third-party integrations (PSP, BVN/NIN verification, SMS/OTP gateways) that touch customer data. It applies to all employees, contractors, and founders, and takes effect the moment any team member suspects — not just confirms — a security incident.

1.3 Out of Scope

- Routine IT support issues with no security implication (e.g., forgotten passwords with no suspicious login pattern).
- Physical security incidents unrelated to information systems (handled under a separate physical security policy).

Definition of a security incident: any event that compromises, or credibly threatens to compromise, the confidentiality, integrity, or availability of PayGuard systems or customer data. This includes attempted intrusions, not only successful ones — CBN's 24-hour reporting rule explicitly covers unsuccessful attacks against a regulated fintech.

2. Roles and Responsibilities

A five-person startup cannot staff a 24-person incident response team, so roles here are function-based, not headcount-based — one person may hold two roles simultaneously, except that the Incident Commander should never also be the Scribe (decision-making and logging pull attention in different directions).

Role	Typical Owner	Responsibility
Incident Commander (IC)	CTO / Head of Engineering	Owns the overall response. Declares severity, authorizes containment actions (e.g., taking a

Role	Typical Owner	Responsibility
		service offline), and is the single point of decision-making during the incident.
Technical Lead	Senior Developer / SysAdmin on duty	Investigates the technical scope, executes containment and eradication steps, coordinates with the IC on system changes.
Communications Lead	Founder / COO / designated comms contact	Owns all outward-facing messaging — customers, press, social media — and internal status updates to the wider team.
Legal & Compliance Lead	External counsel or in-house Compliance Officer	Determines NDPC / CBN notification obligations, drafts regulatory submissions, advises on customer notification wording and liability exposure.
Scribe	Any available team member, assigned at incident start	Maintains the timestamped incident log from the first minute — every action taken, every decision made, every escalation. This log becomes the backbone of the assessment report.
On-call Engineer	Rotating duty roster	First responder outside business hours. Triage the initial alert and escalates to the IC if severity warrants.

2.1 Escalation Rule

Any team member who suspects an incident notifies the on-call engineer or Technical Lead immediately, by phone call — not Slack, not email — since primary tools may be affected by the incident itself. The on-call engineer escalates to the Incident Commander within 15 minutes of confirming the issue is security-relevant.

2.2 Backup Contacts

Every role above must have a named backup, because incidents do not wait for people to be reachable. Maintain the current on-call roster and backup contact list in the Appendix and keep it updated whenever team composition changes.

3. Severity Classification

Severity determines who gets pulled in and how fast the clock starts on regulatory notification. When in doubt, classify one level higher — the cost of over-reacting to a false alarm is far lower than the cost of a missed CBN deadline.

Level	Criteria	Response	Response Time
SEV-1 — Critical	Active payment fraud, customer funds at risk, confirmed data breach, full system compromise	IC engaged immediately; all hands; regulator notification clock starts now	< 15 min
SEV-2 — High	Suspected unauthorized access, credential stuffing	IC engaged; Technical Lead investigates; Comms Lead on standby	< 1 hour

Level	Criteria	Response	Response Time
	in progress, isolated account takeover		
SEV-3 — Medium	Anomalous but unconfirmed activity, single failed intrusion attempt, non-customer-facing system affected	Technical Lead investigates; IC informed	< 4 hours
SEV-4 — Low	Informational alert, policy violation with no data exposure	Logged and reviewed at next team sync	Next business day

4. Incident Playbooks

These three scenarios were chosen because they represent the incident types most likely to affect an early-stage Nigerian fintech: credential-based account takeover, unauthorized server access, and a suspicious transaction spike. The first two are simulated in this project's VirtualBox lab using Kali/Parrot, Nmap, Wireshark, and OWASP ZAP against isolated targets; the tool-to-playbook mapping is summarized in Section 8.3.

4.1 Playbook A — Account Takeover / Credential Stuffing

An attacker uses leaked or guessed credentials — often reused passwords from unrelated breaches — to log into customer accounts, sometimes at scale via automated tooling. Lab evidence for this playbook is produced with OWASP ZAP (scanning a lab-hosted login page for weak session/authentication controls) and Wireshark (capturing a simulated brute-force/login-flood pattern on the lab network).

Detect

- Spike in failed login attempts from a single IP or IP range — visible in a Wireshark capture as repeated authentication requests in a short window.
- Multiple successful logins to different accounts from the same device fingerprint or IP in a short window.
- Customer reports of unauthorized transactions or profile changes they did not make.
- Logins from geographically implausible locations relative to the customer's history.
- OWASP ZAP scan of the login endpoint flags missing rate limiting, weak session tokens, or absent account-lockout controls before an attack even occurs — the same gaps a real attacker would exploit.

Triage / Confirm Severity

- Confirm whether logins are automated (rate, pattern) vs. a single manual account compromise.
- Check whether any successful login resulted in a fund transfer or profile change — this alone escalates to SEV-1.
- Estimate number of accounts affected to scope the incident.

Contain

1. Force-expire sessions and require password reset for all affected accounts.
2. Temporarily rate-limit or block the source IP range(s) at the WAF/firewall.
3. Freeze outbound transactions on affected accounts pending verification.

4. Enable step-up authentication (OTP/MFA) for the affected account cohort if not already active.

Eradicate

5. Identify and close the entry vector (e.g., missing rate limiting on login endpoint, absent CAPTCHA, no MFA enforcement).
6. Rotate any credentials or API keys that may have been exposed in the same breach dataset.
7. Patch the specific control gap identified — do not just address the symptom.

Recover

8. Restore account access only after identity re-verification (not just password reset).
9. Reconcile any fraudulent transactions; initiate reversal/chargeback process where applicable.
10. Notify affected customers per Section 5.2.

Post-Incident Review

Capture: how many accounts, total funds at risk vs. actually lost, time from first anomalous login to detection, and whether MFA would have prevented it. Feed the answer directly into the hardening recommendations.

4.2 Playbook B — Unauthorized Server / Infrastructure Access

An attacker gains access to a server, database, or admin panel — through an exposed port, an unpatched service, weak SSH credentials, or a misconfigured cloud security group. This is the primary scenario this project's lab exercise simulates: an Nmap scan from a Kali/Parrot attacker VM against an isolated target VM identifies exposed services that map to real hardening gaps, with Wireshark used to confirm what traffic those exposed services actually generate.

Detect

- Unexpected open port or service discovered during routine scanning — this is the finding your Nmap lab evidence should demonstrate.
- Unrecognized SSH sessions, new user accounts, or privilege escalation in system logs.
- Unusual outbound traffic (possible data exfiltration or C2 beaconing) — visible in a Wireshark capture as unexpected destination addresses or protocols from the target VM.
- Intrusion detection / cloud provider security alert.

Triage / Confirm Severity

- Determine what the exposed service/port actually allows an attacker to do — the difference between an informational finding and active compromise.
- Check for evidence of actual access (auth logs, new processes, modified files) vs. a theoretical exposure.
- Identify whether the affected system touches customer data or payment processing — this determines whether CBN's 24-hour clock applies.

Contain

11. Isolate the affected host — remove from load balancer, apply restrictive security group / firewall rule, or disconnect from network entirely if compromise is confirmed.
12. Preserve evidence before making changes: snapshot the disk/VM, export relevant logs, do not reboot a compromised host (this destroys volatile evidence).
13. Revoke and rotate any credentials or keys accessible from the affected host.

Eradicate

14. Close the exposed port/service or restrict it to an internal-only network / VPN.
15. Patch the vulnerable software version identified.
16. Rebuild the host from a known-clean image rather than trying to 'clean' a compromised system.

Recover

17. Bring the rebuilt/patched host back online in stages, monitoring closely for 48–72 hours.
18. Re-run the vulnerability scan (the same Nmap process used in detection) to confirm the finding is closed.
19. Update firewall/security-group baseline documentation to prevent regression.

Post-Incident Review

This is the playbook your assessment report should walk through end-to-end using your lab evidence: the Nmap finding, the Wireshark traffic confirmation, the classification decision, the containment step taken in the lab, and the specific hardening recommendation it produced.

4.3 Playbook C — Payment Fraud / Suspicious Transaction Spike

A sudden spike in transaction volume, value, or a pattern (e.g., many small transactions just under a fraud-detection threshold) suggests fraud — either external attack or an internal control failure.

Detect

- Transaction monitoring system flags volume or value anomaly against baseline.
- Multiple transactions just under a known approval/review threshold (structuring pattern).
- Spike in transactions to a small number of destination accounts.
- Customer complaints of unauthorized transactions arriving faster than the normal support volume.

Triage / Confirm Severity

- Quantify exposure: total value at risk, number of accounts and transactions involved, whether transactions are still in-flight or already settled.
- Determine if this correlates with a known account-takeover incident (Playbook A) or is a distinct fraud pattern.

Contain

20. Pause outbound settlement/disbursement for the affected transaction batch if still in-flight.
21. Apply a temporary hold or step-up review on the destination accounts receiving the suspicious inflow.
22. Notify the payment processor / PSP partner to flag or reverse transactions where their systems allow it.

Eradicate

23. Identify the control gap that allowed the pattern through (e.g., threshold set too high, missing velocity checks, absent destination-account risk scoring).
24. Adjust fraud-rule thresholds and velocity checks; deploy the fix and validate against historical data.

Recover

25. Reconcile ledgers; process reversals/chargebacks; document any funds that cannot be recovered.
26. Resume normal transaction processing once monitoring confirms the anomaly has stopped.

Post-Incident Review

This playbook has the most direct CBN exposure — any confirmed fraud affecting customer funds should be treated as SEV-1 and put you inside the 24-hour notification clock immediately.

5. Communications Plan

5.1 Internal Communications

Primary channel: a dedicated incident Slack/WhatsApp channel created the moment an incident is declared SEV-2 or higher. Because incidents can take down primary tooling, maintain a fallback channel (a phone tree or a secondary messaging app on a non-work device) that does not depend on company infrastructure.

- SEV-1: immediate call to the Incident Commander; all relevant roles paged within 15 minutes.
- SEV-2: Incident Commander and Technical Lead notified within 1 hour; Comms Lead placed on standby.
- SEV-3/4: logged and raised at the next scheduled team sync; no immediate paging required.

The Scribe posts a status update to the incident channel at least every 30 minutes during a SEV-1, even if the update is simply "still investigating, no new information." Silence during an active incident is what causes panic to spread through the rest of the team.

5.2 Customer Communications

Two templates, used depending on how much is confirmed at the time communication is required:

Holding statement (used when facts are still being confirmed)

"We are aware of unusual activity affecting some accounts and are actively investigating. As a precaution, [specific protective action, e.g. 'we have temporarily paused transfers on affected accounts']. We will update you within [X hours] with more information. Your funds are protected under [relevant guarantee/policy]."

Full disclosure (used once facts are confirmed)

"On [date], we identified [what happened] affecting [scope]. [What data/funds were involved]. We have [containment/remediation actions taken]. [What the customer should do, if anything]. We have notified the appropriate regulators. If you have questions, contact us at [channel]."

Never speculate about root cause or attribute blame in a customer-facing statement before the Technical Lead and Legal & Compliance Lead have both signed off — an inaccurate early statement is harder to walk back than a short delay.

5.3 Regulatory Communications

The table below reflects current Nigerian requirements as of August 2026. Because CBN's window (24 hours) is shorter than NDPC's (72 hours), PayGuard's operative internal deadline for any incident touching customer data is 24 hours — the shorter of the two always governs, the same logic Nigerian regulators apply when frameworks overlap.

Regulator / Party	Deadline	Required Content	Owner
NDPC	Within 72 hours of becoming aware of a breach likely	Nature of the breach; categories and approximate number of data subjects	Legal & Compliance Lead

Regulator / Party	Deadline	Required Content	Owner
	to pose a risk to data subjects' rights or freedoms (NDPA 2023, Section 40)	affected; likely consequences; measures taken or proposed	
CBN (Banking Supervision Dept.)	Cyber incidents must be reported to the CBN within 24 hours. Regular audits, metrics reviews, and compliance checks are mandatory. The CBN enforces adherence through spot checks and sector-wide audits	Incident description, systems affected, initial impact assessment; follow-up report once investigation completes	Legal & Compliance Lead + IC
Affected Data Subjects (customers)	"Without undue delay" once high risk to their rights is confirmed	What happened, what data was affected, what they should do, what PayGuard is doing	Communications Lead

6. Reporting

Every incident — regardless of severity — produces two artifacts: a live timeline maintained during the incident by the Scribe, and a post-incident report finalized within 5 business days of resolution. The timeline below is the template your assessment report's evidence trail should follow.

6.1 Incident Timeline Template

Timestamp	Event	Detail
[HH:MM]	Detection	e.g., 'Nmap scan flagged port 3306 (MySQL) exposed on public interface of app-server-02'
[HH:MM]	Escalation	e.g., 'On-call engineer notified Technical Lead by phone'
[HH:MM]	Severity assigned	e.g., 'Classified SEV-2 — customer data on affected DB but no confirmed access'
[HH:MM]	Containment action	e.g., 'Security group rule applied restricting port 3306 to internal VPC only'
[HH:MM]	Evidence preserved	e.g., 'Disk snapshot taken; auth logs exported to secure storage before any remediation'
[HH:MM]	Regulatory clock started	e.g., 'CBN 24-hour window begins; Legal & Compliance Lead notified'

6.2 Post-Incident Report Structure

27. Executive summary — one paragraph, written for a non-technical founder or board member.
28. Incident classification and timeline — pulled directly from the Scribe's log.
29. Root cause — the specific technical or process gap that allowed the incident, not just the symptom.
30. Impact assessment — accounts/data/funds affected, actual vs. potential.
31. Response actions taken — containment, eradication, recovery steps.
32. Regulatory notifications sent — what, when, to whom, and confirmation of receipt where available.
33. Remediation and hardening recommendations — specific, owned, and dated (see Section 7).
34. Lessons learned — what would have caught this sooner, and what in this plan needs to change.

7. Remediation & Hardening Recommendations

Recommendations should trace directly back to a specific finding — from a lab scan, a real incident, or a control gap identified during playbook testing. Generic 'improve security' recommendations do not belong in this table.

Recommendation	Priority	Owner	Status	Target Date
Enforce MFA on all customer-facing login endpoints	High	Eng	Open	—
Add rate limiting / CAPTCHA on login endpoint	High	Eng	Open	—
Close non-essential open ports identified in scan (see lab evidence)	High	Eng	Open	—
Patch identified service versions to current stable release	Medium	Eng	Open	—
Implement velocity/threshold checks on transaction monitoring	High	Eng	Open	—
Formalize on-call roster and backup contacts	Medium	Ops	Open	—

8. Appendix

8.1 Contact & Escalation Roster

Role	Name	Phone	Email
Incident Commander	[Name]	[Phone]	[Email]
Technical Lead	[Name]	[Phone]	[Email]
Communications Lead	[Name]	[Phone]	[Email]
Legal & Compliance Lead	[Name]	[Phone]	[Email]

Role	Name	Phone	Email
On-call Engineer (rotation)	[Name / Roster link]	[Phone]	[Email]

8.2 Lab Environment Notes (for Assessment Evidence)

All offensive and defensive testing referenced in this plan (Nmap scans, Wireshark captures, OWASP ZAP scans, Kali/Parrot-based reconnaissance) was performed exclusively within an isolated VirtualBox lab environment, on a host-only or internal network with no route to production systems or the public internet. No scanning, capture, or exploitation activity described in this document was performed against any live, third-party, or production system. Screenshots, packet captures, and scan reports from the lab are retained as supporting evidence in the accompanying assessment report.

8.3 Tools Reference

Tool	Role in Lab	How It Is Used
VirtualBox	Lab hypervisor	Hosts the isolated target VM(s) and the Kali/Parrot attacker VM on a host-only or internal-only virtual network with no route to production or the public internet.
Kali Linux / Parrot OS	Attacker / testing VM	Source platform for reconnaissance, scanning, and traffic capture — all tools below run from this VM.
Nmap	Network & service discovery	Identifies open ports, running services, and software versions on the lab target — primary evidence source for Playbook B (Unauthorized Server Access).
Wireshark	Network traffic capture & analysis	Captures and inspects live traffic on the lab network — used to demonstrate detection of anomalous patterns such as a login-flood/brute-force attempt (Playbook A) or unusual outbound traffic (Playbook B).
OWASP ZAP	Web application vulnerability scanning	Scans a lab-hosted web application/API for authentication, session-management, and injection weaknesses — primary evidence source for Playbook A (Account Takeover / Credential Stuffing).

8.4 Revision History

v1.0 — August 2026 — Initial capstone draft.

v1.1 — August 2026 — Added glossary, references, and expanded lab tooling (Wireshark, OWASP ZAP, Parrot OS) per instructor-suggested tool list; tool-to-playbook mapping added to Section 4.

9. Glossary

Terms are defined as used within this plan; some (e.g., MFA, CVE) have broader technical meanings not fully captured here.

Term	Definition
Account Takeover (ATO)	An attack where a threat actor gains unauthorized control of a legitimate user's account, typically using stolen, leaked, or guessed credentials.
Active Scan	In OWASP ZAP, a scan mode that actively sends crafted requests to a target application to test for vulnerabilities (e.g., injection, XSS), as opposed to a passive scan that only observes existing traffic.
Anti-Clickjacking Header	An HTTP response header (such as X-Frame-Options or a CSP frame-ancestors directive) that prevents a page from being loaded inside a hidden frame on a malicious site — its absence allows clickjacking attacks.
BVN / NIN	Bank Verification Number / National Identification Number — identity verification numbers used across Nigerian financial services.
C2 (Command & Control)	Infrastructure an attacker uses to communicate with and control compromised systems after a breach.
CBN	Central Bank of Nigeria — the apex regulator for Nigerian banks, payment service providers (PSPs), and other licensed financial institutions.
Containment	The incident response phase focused on stopping an incident from spreading or causing further damage, before the root cause is fully removed.
Content Security Policy (CSP)	An HTTP response header that lets a site restrict which sources of scripts, styles, and other content a browser is allowed to load, reducing the impact of XSS and injection attacks.
Cross-Site Request Forgery (CSRF)	An attack that tricks a logged-in user's browser into submitting an unwanted request to a site they're authenticated on, exploiting the site's trust in the user rather than the user's trust in the site.
Cross-Site Scripting (XSS)	A vulnerability class where an attacker injects malicious script into a web page viewed by other users, exploiting the user's trust in the site.
CVE	Common Vulnerabilities and Exposures — a public, standardized identifier for a known software vulnerability.
Data Controller / Processor	Under the NDPA, the entity that determines the purpose of processing personal data (controller) versus the entity that processes it on the controller's behalf (processor).
Data Subject	An identifiable individual whose personal data is collected or processed — i.e., a PayGuard customer.
DVWA (Damn Vulnerable Web App)	An intentionally vulnerable PHP/MySQL web application, bundled with Metasploitable2, used for practicing and demonstrating web application security testing in a legal, controlled environment.
Eradication	The incident response phase focused on removing the root cause of an incident (e.g., patching a vulnerability, rotating compromised credentials).
Exfiltration	The unauthorized transfer of data out of an organization's systems, typically by an attacker.

Term	Definition
HttpOnly Flag	A cookie attribute that prevents client-side JavaScript from reading the cookie's value, reducing the risk of session theft via XSS.
MFA / OTP	Multi-Factor Authentication / One-Time Password — authentication methods requiring more than a password to verify identity.
NDPA	Nigeria Data Protection Act, 2023 — Nigeria's principal data protection law.
NDPC	Nigeria Data Protection Commission — the regulator established under the NDPA to enforce data protection obligations.
Nmap	Network Mapper — an open-source tool used to discover hosts, open ports, and running services on a network, and to fingerprint software versions.
OWASP ZAP	Zed Attack Proxy — an open-source web application security scanner used to find vulnerabilities such as injection flaws, broken authentication, and session management issues.
Playbook	A predefined, step-by-step response procedure for a specific type of incident.
PSP	Payment Service Provider — a CBN-licensed entity that facilitates electronic payment transactions.
RACI	Responsible, Accountable, Consulted, informed — a framework for clarifying roles and decision rights.
Remote Code Execution (RCE)	A vulnerability that allows an attacker to run arbitrary commands or code on a target system, typically the most severe class of application or service vulnerability.
SameSite Attribute	A cookie attribute that restricts whether a cookie is sent along with cross-site requests, helping to mitigate CSRF and related attacks.
Severity (SEV) Level	A classification (SEV-1 to SEV-4 in this plan) indicating the urgency and impact of an incident, used to determine response speed and escalation.
Spider (Web Crawler)	In OWASP ZAP, the component that automatically discovers pages and links within a target application by following URLs, building a map of the site prior to active scanning.
Threat Intelligence	Information about existing or emerging threats (e.g., known malicious IPs, leaked credential dumps) used to inform detection and defense.
Velocity Check	A fraud-detection control that flags transactions occurring at an unusually high rate or frequency relative to a customer's normal behavior.
Wireshark	An open-source network protocol analyzer used to capture and inspect live network traffic in detail.

10. References

Regulatory sources:

- Nigeria Data Protection Act, 2023, Section 40 (breach notification obligations). Nigeria Data Protection Commission — ndpc.gov.ng.
- Central Bank of Nigeria, Risk-Based Cybersecurity Framework and Guidelines for Deposit Money Banks and Payment Service Providers (effective 1 January 2019; updated 2024) — cbn.gov.ng.
- Central Bank of Nigeria, Banks and Other Financial Institutions Act (BOFIA) 2020 — statutory basis for CBN's cybersecurity supervisory powers.

Secondary/supporting sources consulted for regulatory timeline summaries (verify against primary regulator text before external use):

- "Nigeria's Data Protection Act: What Businesses Should Know to Comply" — Pandects, 2026.
- "Nigeria Data Protection Act (NDPA) 2023: Complete Compliance Checklist" — Dimeri, 2026.
- "Secure Banking Solutions in Nigeria" — Simpa Labs, 2026 (CBN 24-hour breach reporting and control-requirement summary).
- "Cyber-Security Regulation in the Nigerian FinTech Industry" — The New Practice (TNP), legal analysis of the CBN Risk-Based Cybersecurity Framework.

Frameworks referenced for general incident-response structure (Detect–Triage–Contain–Eradicate–Recover–Review):

- NIST SP 800-61 Rev. 2, Computer Security Incident Handling Guide — used as a general structural reference, not a compliance requirement.

10.1 Note on Citation Practice

This plan is a coursework capstone artifact. Where regulatory figures (e.g., the 24-hour and 72-hour notification windows) directly affect compliance decisions in a real deployment, confirm current wording against the primary source — cbn.gov.ng and ndpc.gov.ng — since regulatory circulars are amended periodically and secondary summaries can lag behind the current text.